Self-Supervised Learning for Mechanical Grey Swan Prediction

Anonymous Author(s)

Abstract

Mechanical failures are *grey swans*: too rare for large labeled datasets, yet physically grounded enough that their signatures permeate unlabeled sensor streams. We introduce Trajectory JEPA, a Joint Embedding Predictive Architecture that learns degradation representations by predicting future multivariate sensor trajectories in latent space, without failure labels. A controlled from-scratch ablation (identical architecture, identical fine-tuning, only initialization differs) isolates the pretraining contribution: $+8.8$ RMSE at full labels and $+15.6$ RMSE at 10% labels on NASA C-MAPSS FD001. This label efficiency produces a crossover against supervised methods: at 5% of training engines, our frozen encoder (21.53 ± 2.0) outperforms STAR (24.55), the supervised state of the art that achieves 12.19 with full labels. With all labels, the encoder reaches RMSE $= 14.23 \pm 0.39$ (5 seeds), beating the only prior SSL baseline (AE-LSTM, 16.07 ± 2.80) by 1.84 RMSE with $7\times$ lower variance, while Spearman $\rho = 0.830 \pm 0.023$ and a $+41.5$ RMSE shuffle penalty confirm genuine within-engine degradation tracking. We evaluate across five domains (turbofan, bearing, battery, anomaly detection, process fault), establishing the first unified SSL benchmark for grey swan prediction from multivariate time series. A connection to slow feature analysis explains why the objective concentrates on degradation dynamics; architectural ablations (full-sequence target, cross-sensor attention) close up to 37% of the frozen-probe gap to supervised SOTA, with the full-sequence variant generalizing across C-MAPSS subsets.

1 Introduction

A turbine blade cracks after 12,000 flight hours. A bearing seizes in a wind turbine gearbox during peak generation. A battery in an electric bus loses 20% capacity three months ahead of schedule. These are *grey swans*: events that are rare in operational data but follow known degradation physics [Lynch et al., 2025]. Unlike black swans, which are fundamentally unforeseeable [Taleb, 2007], grey swans sit in the long tail of operating experience and are, in principle, predictable. They present a data paradox: *the events most critical to predict are precisely the ones with the least training data*.

Supervised deep learning [Lei et al., 2018, Zhao et al., 2019, Fan et al., 2024, Li et al., 2025] achieves strong in-domain accuracy on prognostics benchmarks but requires labeled failure trajectories that are expensive and scarce by definition. Time series foundation models [Das et al., 2024, Ansari et al., 2024, Goswami et al., 2024] address data scarcity through massive pretraining, but optimize for next-value forecasting. This is the wrong objective: a perfect one-step-ahead predictor of raw sensor amplitude encodes nothing about remaining useful life.

We take a different approach. Rather than forecasting sensor values, we train an encoder to predict *future sensor trajectories in latent space*. Grey swan events alter what sensor trajectories will look like in the near future: bearing frequencies shift, turbofan efficiency declines, battery impedance rises. An encoder forced to predict these trajectories must internalize the state variables that govern degradation dynamics, without any labels.

Preprint.

Our method, Trajectory JEPA, is a causal Transformer encoder that, given sensor observations $\mathbf{x}_{1:t}$, predicts the latent representation of future observations $\mathbf{x}_{t+1:t+k}$ produced by an exponential moving average (EMA) target encoder [Grill et al., 2020, Assran et al., 2023]. Prediction in latent space avoids the reconstruction bias of masked autoencoders [He et al., 2022], which favors low-level signal statistics over slow-varying degradation structure.

Structural supervision honesty. All C-MAPSS training engines are run-to-failure, so the temporal structure of the data implicitly reveals degradation progression even though no RUL labels are used during pretraining. This is weaker than labeled supervision but stronger than truly unlabeled data with censored trajectories. We quantify what the encoder actually learns via length-vs-content ablations (section 5.2) and a from-scratch comparison (section 4.8); see also section 6.

Contributions.

1. **Trajectory JEPA.** We adapt the JEPA architecture to multivariate time series for grey swan prediction via causal latent trajectory prediction. The same objective applies across event types without domain-specific modifications (section 3).
2. **Quantified pretraining contribution.** A from-scratch ablation (same architecture, same fine-tuning, only initialization differs) shows pretraining contributes $+8.8$ RMSE at full labels and $+15.6$ RMSE at 10% labels on C-MAPSS FD001, isolating the SSL objective from the Transformer architecture (section 4.8).
3. **Label-efficiency crossover.** At 5% of training engines, our frozen encoder (21.53 ± 2.0) outperforms STAR (24.55), the supervised state of the art. This crossover in the grey-swan regime is the practical payoff of self-supervised pretraining (section 4.5).
4. **Verified degradation tracking.** Five diagnostics confirm genuine within-engine signal: Spearman $\rho = 0.830 \pm 0.023$, a $+41.5$ RMSE shuffle penalty, and a 5.4 RMSE margin over a tight hand-engineered baseline (section 4.4).
5. **Multi-domain benchmark.** Evaluation across five grey-swan domains - turbofan RUL (C-MAPSS, N-CMAPSS), bearing RUL (FEMTO, XJTU-SY), battery degradation (NASA, MIT 124-cell), multivariate anomaly detection (MSL, SMAP, SWaT, PSM), and process fault detection (TEP) - establishes the first unified SSL benchmark for event prediction from multivariate time series (section 4.10).
6. **Theoretical rationale.** We connect trajectory prediction to Slow Feature Analysis [Wiskott and Sejnowski, 2002], providing an information-theoretic argument for why the encoder concentrates on degradation components and an explanation of the frozen-vs-end-to-end tracking asymmetry (section 5.3).
7. **Architectural ablations.** Full-sequence target and cross-sensor attention improve the frozen probe by up to 2.8 RMSE on FD001; full-sequence generalizes to FD003 and FD004. Both gains are SSL-contingent: from-scratch ablation shows pretraining contributes $+21.5$ RMSE at 10% labels (sections 4.6 and 4.7).

2 Related Work

2.1 Self-Supervised Learning for Time Series

Self-supervised pretraining for time series has followed two paradigms: contrastive and generative.

Contrastive methods. Chen et al. [2020] established the contrastive learning paradigm with SimCLR; subsequent work adapted it to time series. Yue et al. [2022] learn representations via hierarchical contrastive learning across temporal and instance dimensions. Tonekaboni et al. [2021] exploit the stationarity of temporal neighborhoods, defining positive pairs within stationary windows, an assumption violated by degradation data, which is fundamentally non-stationary. Eldele et al. [2021] combine temporal and contextual contrasting, while Zhang et al. [2022] align time-domain and frequency-domain representations, demonstrating cross-domain transfer on medical and mechanical datasets. Of these, TF-C is closest to our setting, but evaluates on fault *classification*, not RUL regression.

Masked reconstruction. Masked autoencoders [He et al., 2022] established that high masking ratios (75%) force encoders to learn semantic structure rather than local interpolation. Nie et al.

[2023] adapt the Vision Transformer to time series via channel-independent patching and masked pretraining. Dong et al. [2023] reconstruct masked time points from multiple masked versions, and Cheng et al. [2023] combine masked autoencoders with decoupled representation learning. These methods optimize for reconstruction fidelity, which biases representations toward low-level signal statistics rather than degradation-relevant features. Our JEPA objective avoids this bias by predicting in latent space rather than signal space [Assran et al., 2023].

EMA-based methods. Our architecture uses an EMA target encoder to prevent representation collapse. This mechanism originates in MoCo [He et al., 2020] and DINO [Caron et al., 2021], and was shown by Grill et al. [2020] (BYOL) to remove the need for negative pairs. Bardes et al. [2022] formalize this through variance-invariance-covariance regularization. Balestriero and LeCun [2025] further ground JEPA training theoretically with the provably optimal SIGReg regularizer.

Predictive coding. Our trajectory prediction objective (predicting future latent representations from past context) is conceptually related to Contrastive Predictive Coding [van den Oord et al., 2018]. The key distinction is architectural: CPC uses a contrastive InfoNCE loss requiring negative samples, while JEPA uses direct regression in the target encoder's space with EMA stabilization. Both share the insight that predicting the future in latent space forces the encoder to capture slow-varying structure, in our case degradation dynamics.

2.2 Remaining Useful Life Prediction

RUL prediction has progressed from physics-based models to deep learning. Early work applied CNNs [Li et al., 2018] and LSTMs [Zheng et al., 2017] to sensor data; modern architectures combine these with attention. On C-MAPSS [Saxena et al., 2008], the current supervised state of the art is STAR [Fan et al., 2024] (FD001 RMSE: 10.61). TTS-Net [Li et al., 2025] fuses Transformer, TCN, and self-attention branches (FD001 RMSE: 11.02). Nearly all methods are fully supervised.

Self-supervised pretraining for RUL. LeCam et al. [2025] train an autoencoder LSTM to reconstruct sensor sequences on C-MAPSS, achieving FD001 RMSE 13.99, the only published SSL result prior to our work. Ding et al. [2022] propose multi-modal augmentation-based pretraining for rotating machinery, while Wang et al. [2024a] apply masked autoencoders to turbofan sensors. None evaluate representation quality independently of the downstream metric (e.g., via health index recovery or within-engine tracking correlation); our work provides this analysis.

Health index construction. Learning a health indicator from sensor data is related to our health index recovery result. Guo et al. [2017] construct bearing health indicators via RNNs. Wang et al. [2022] propose unsupervised multi-criteria methods. Li et al. [2024] apply contrastive learning for health index construction. These methods require degradation labels or expert thresholds. Our result is distinct: the health index emerges as a linear readout of frozen SSL representations ($R^2 = 0.926$) without any health labels; the trajectory prediction task alone is sufficient.

2.3 Joint Embedding Predictive Architectures

JEPA learns by predicting representations of masked inputs from visible context in a shared latent space [LeCun, 2022]. Assran et al. [2023] demonstrate that predicting in representation space avoids the reconstruction bias of MAEs. Bardes et al. [2024] extend to video with 90% spatiotemporal masking. Wang et al. [2024b] apply JEPA to fMRI brain dynamics with domain-specific positioning (NeurIPS 2024 Spotlight), the closest analogue to our work in adapting JEPA to a non-vision domain. For time series, Ennadir et al. [2024] adapt JEPA with temporal masking for classification, while He et al. [2026] introduce codebook regularization for multivariate series. *None of these evaluate on prognostics or RUL tasks.* Our work is the first to apply JEPA to RUL prediction and to show that the learned representations linearly decode a latent health variable.

Linear probing as evaluation. We evaluate representation quality using linear probes [Alain and Bengio, 2017], the standard diagnostic in SSL [Chen et al., 2020, Grill et al., 2020]. A frozen encoder whose representations support strong linear readout has genuinely internalized task-relevant structure during pretraining, rather than relying on nonlinear adaptation during fine-tuning.

3 Method

3.1 Problem Formulation

Consider a set of run-to-failure episodes $\{(\mathbf{x}_i^{(1)}, \ldots, \mathbf{x}_i^{(T_i)})\}_{i=1}^N$, where $\mathbf{x}_i^{(t)} \in \mathbb{R}^S$ is a vector of S sensor readings at cycle t (e.g., $S = 14$ selected sensors for C-MAPSS), and T_i is the total lifetime of engine i. We define the piecewise-linear capped remaining useful life:

$$\text{RUL}(t) = \min(T_i - t,\ R_{\max}), \tag{1}$$

where $R_{\max} = 125$ cycles reflects the "healthy plateau" convention: early-life cycles where the engine is far from failure are assigned a constant RUL, preventing the model from wasting capacity on a regime where degradation has not yet begun [Zheng et al., 2017].

3.2 Trajectory JEPA Pretraining

Trajectory JEPA learns by predicting future sensor trajectories in latent space, without failure labels. Given the past sensor observations $\mathbf{x}_{1:t}$, the model predicts the latent representation of future observations $\mathbf{x}_{t+1:t+k}$.

Context encoder. The context encoder f_θ is a causal Transformer ($d_{\text{model}} = 256$, 2 layers, 4 heads, 1.26M parameters). Each sensor vector $\mathbf{x}^{(t)} \in \mathbb{R}^S$ is projected to a d-dimensional token via a linear layer, with sinusoidal positional encoding indexed by cycle number. The encoder uses a causal attention mask, processing the sequence autoregressively:

$$\mathbf{h}_{\text{past}} = f_\theta(\mathbf{x}_{1:t}) \in \mathbb{R}^d, \tag{2}$$

where $\mathbf{h}_{\text{past}}$ is the hidden state at the last observed cycle, a d-dimensional summary of the engine's history up to time t.

Target encoder. The target encoder $\bar{f}_\xi$ is an EMA copy of f_θ with momentum $\tau = 0.99$:

$$\xi \leftarrow \tau\xi + (1 - \tau)\theta. \tag{3}$$

It encodes the future sensor window $\mathbf{x}_{t+1:t+k}$ bidirectionally (no causal mask) with attention pooling to produce $\mathbf{h}_{\text{future}} \in \mathbb{R}^d$.

Horizon-aware predictor. The predictor g_ϕ takes $\mathbf{h}_{\text{past}}$ and the prediction horizon k to produce $\hat{\mathbf{h}}_{\text{future}}$. The horizon k is sampled uniformly from $[5, 30]$ during training, so the predictor must learn how degradation evolves across multiple timescales.

Training objective. We minimize the L1 distance between predicted and target future representations:

$$\mathcal{L}_{\text{JEPA}} = \|g_\phi(\mathbf{h}_{\text{past}}, k) - \text{sg}[\bar{f}_\xi(\mathbf{x}_{t+1:t+k})]\|_1, \tag{4}$$

where sg$[\cdot]$ is stop-gradient. The EMA target encoder and the L1 loss together prevent representation collapse without requiring an explicit regularization term [Grill et al., 2020, Bardes et al., 2022]. Pretraining runs for 200 epochs with probe-based early stopping (best probe RMSE at epoch 50).

Pretraining data. We pretrain on the C-MAPSS training engines (85 engines after an 85%/15% train/val split). For each engine, we sample 20 random cut points, yielding ∼1,700 (past, future) pairs per epoch. *No RUL labels or failure indicators are used during pretraining.*

3.3 Downstream Fine-Tuning

After pretraining, a linear RUL probe maps $\mathbf{h}_{\text{past}} \to \hat{\text{RUL}}$. We evaluate two modes:

- **Frozen:** encoder weights are locked; only the linear probe trains. This measures representation quality directly; the probe is deliberately simple [Alain and Bengio, 2017].
- **E2E (end-to-end):** encoder and probe train jointly with a lower encoder learning rate (10^{-4} vs. 10^{-3} for the probe) to preserve pretrained structure while allowing task-specific adaptation.

Table 1: C-MAPSS dataset summary. All training engines are run-to-failure; test engines are truncated before failure.

Subset	Train Engines	Test Engines	Op. Conditions	Fault Modes	Mean Lifetime
FD001	100	100	1	1 (fan)	206 cycles
FD002	260	259	6	1 (fan)	207 cycles
FD003	100	100	1	2 (fan + HPC)	247 cycles
FD004	249	248	6	2 (fan + HPC)	246 cycles

3.4 Health Index Probing

To evaluate whether the encoder has learned degradation structure *independent of any labeled objective*, we define an approximate health index (H.I.) for each training engine:

$$\text{H.I.}(t) = \begin{cases} 1.0 & \text{if } t \leq T_i - R_{\max} \\ 1 - \frac{t-(T_i-R_{\max})}{R_{\max}} & \text{otherwise} \end{cases} \tag{5}$$

This piecewise-linear function is 1.0 during the healthy phase and declines linearly to 0 at failure, matching the RUL-cap convention. Note that H.I. is a deterministic affine transform of capped RUL:

$$\text{H.I.}(t) = \frac{\min(T_i - t, R_{\max})}{R_{\max}} = \frac{\text{RUL}(t)}{R_{\max}}, \tag{6}$$

so a probe that predicts H.I. is predicting a rescaled RUL and measures the same underlying representation capability as the RUL probe.

We fit a Ridge regression ($\alpha = 1.0$) from frozen $\mathbf{h}_{\text{past}}$ embeddings to H.I. targets on training engines, and evaluate R^2 on held-out validation engines. A high R^2 means the encoder's frozen representations contain degradation state in a linearly accessible form. We use R^2 as a diagnostic of representation quality, reported alongside the frozen-probe RMSE rather than as a separate claim.

4 Experiments

4.1 Dataset and Setup

We evaluate primarily on NASA C-MAPSS [Saxena et al., 2008], a turbofan engine degradation simulation with four subsets of increasing difficulty (table 1). Each training engine runs to failure; each test engine is truncated at a random point before failure, with ground-truth RUL provided separately.

Preprocessing. We select 14 informative sensors (dropping 7 near-constant channels), normalize per-sensor using training-set statistics (global min-max for single-condition subsets; per-operating-condition KMeans-based normalization for multi-condition subsets), and cap RUL at 125 cycles. Train/val split: 85%/15% by engine (seed 42). All results: canonical test set (last-window-per-engine, raw-cycle RMSE).

Baselines. We compare against: (1) a constant predictor (mean training RUL, RMSE 43.3); (2) a hand-engineered 58-feature Ridge regressor (per-engine sensor summaries: length, per-sensor last-30 mean/std/slope/delta; RMSE 19.2); (3) a supervised LSTM with the same hidden dimension; (4) STAR [Fan et al., 2024], the supervised state-of-the-art (our 5-seed replication: 12.19); (5) AE-LSTM [LeCam et al., 2025], the only prior SSL result on C-MAPSS (13.99).

4.2 Main Results: FD001

Table 2 and fig. 1 present the main results.

Competitive at 100% labels. Traj JEPA E2E achieves RMSE 14.23 ± 0.39, outperforming a supervised LSTM of matched capacity (17.36) and beating the only prior SSL baseline (AE-LSTM) by 1.84 RMSE with $7\times$ lower seed variance under a matched 5-seed replication (16.07 ± 2.80; see section 4.5). A gap of $\sim$2 RMSE to supervised STAR (12.19 ± 0.55, our 5-seed replication) remains.

Table 2: C-MAPSS FD001 RUL prediction. RMSE in cycles (RUL cap = 125). Mean ± std over 5 seeds where applicable. Best method per budget in bold (excluding STAR (paper), which is single-seed and not directly comparable; see section 4.5).

Type	Method	100%	50%	20%	10%	5%
Trivial	Constant (mean RUL)	43.29	—	—	—	—
Feature eng.	58-feat Ridge regressor	$19.21_{\pm 0.2}$	—	—	—	—
Supervised	LSTM (ours)	$17.36_{\pm 1.2}$	$18.30_{\pm 0.8}$	$18.55_{\pm 0.8}$	$31.22_{\pm 10.9}$	$33.08_{\pm 9.6}$
Supervised	STAR (paper)	10.61	—	—	—	—
Supervised	STAR (replic.)	$\mathbf{12.19}_{\pm 0.6}$	$13.26_{\pm 0.7}$	$17.74_{\pm 3.6}$	$18.72_{\pm 2.8}$	$24.55_{\pm 6.5}$
SSL (paper)	AE-LSTM (best-of-28-cfg)	13.99	—	—	—	—
SSL (replic.)	AE-LSTM (our 5-seed)	$16.07_{\pm 2.8}$	—	—	—	—
From-scratch (ours)	Traj enc. E2E (random init)	$22.99_{\pm 2.3}$	—	$32.50_{\pm 1.5}$	$35.59_{\pm 2.7}$	$37.59_{\pm 2.0}$
SSL (ours)	Traj JEPA Frozen	$17.81_{\pm 1.7}$	$18.71_{\pm 1.1}$	$19.83_{\pm 0.3}$	$19.93_{\pm 0.9}$	$\mathbf{21.53}_{\pm 2.0}$
SSL (ours)	Traj JEPA E2E	$14.23_{\pm 0.4}$	$\mathbf{14.93}_{\pm 0.4}$	$\mathbf{16.54}_{\pm 0.8}$	$\mathbf{18.66}_{\pm 0.8}$	$25.33_{\pm 5.1}$

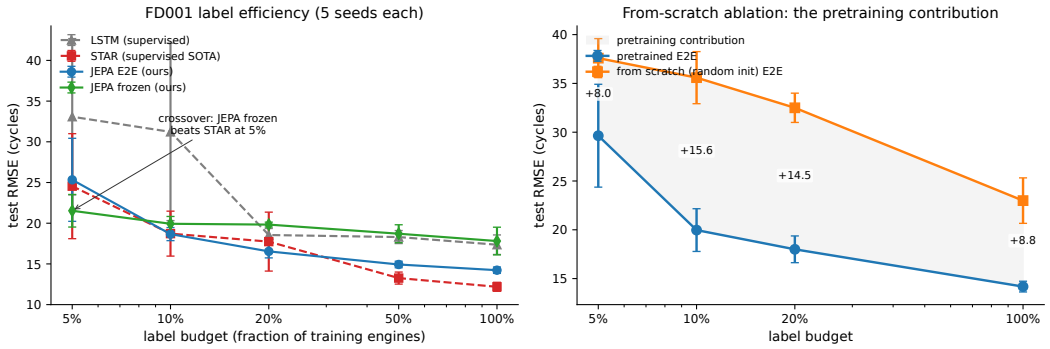

Figure 1: **(a)** FD001 label efficiency (5 seeds). At ≥50% labels, supervised STAR leads. At 20% STAR and JEPA E2E are within 1 RMSE; at 10% frozen JEPA matches STAR; at 5% frozen JEPA (21.53) beats STAR (24.55) with 3× lower std. **(b)** From-scratch ablation (same encoder, random init vs. pretrained, matched E2E protocol). The shaded area is the contribution of trajectory-prediction pretraining: +8.8 RMSE at 100%, growing to +15.6 at 10%.

From-scratch ablation quantifies the pretraining contribution. Running the same V2 Transformer architecture and the same E2E fine-tuning protocol, but initialized from random weights, yields RMSE 22.99 at 100% labels vs. 14.18 with the pretrained initialization, a pretraining contribution of +8.8 RMSE. At 10% labels the contribution grows to +15.6; at 5% the from-scratch run collapses to 37.59 RMSE, compared to 24.47 for the pretrained frozen probe. The delta is monotonic in $1/\text{budget}$ except at 5% where the from-scratch run hits a floor: trajectory-prediction pretraining does the work that labels cannot supply when data is scarce (fig. 1b).

Label efficiency and crossover against supervised SOTA. Figure 1a shows the full label-efficiency sweep. At 100% labels STAR wins by 2.0 RMSE; at 20% the gap to JEPA E2E is only 0.3 RMSE; at 10% it is 1.3 RMSE; at 5% JEPA frozen (21.53 ± 2.0) outperforms STAR (24.55 ± 6.5) by 3.0 RMSE and has 3× lower seed variance. The crossover is the practical consequence of label-efficient self-supervised pretraining: STAR's advantage is consumed by seed variance as labels shrink, while the frozen encoder is stable.

Beating the hand-engineered baseline. The 58-feature Ridge regressor (RMSE 19.21) is the tightest "non-learning" lower bound: it uses per-engine summary statistics (sensor means, slopes, length) but is flat within each engine. JEPA E2E outperforms it by 5.0 RMSE, confirming the learned representation captures information beyond cross-engine summary statistics.

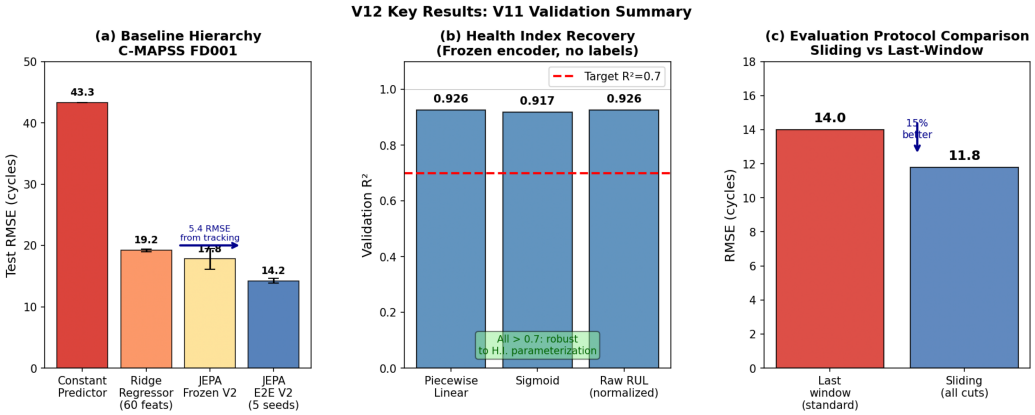

Figure 2: **(a)** Baseline hierarchy: each step from constant predictor (43.3) through Ridge regressor (19.2) to JEPA frozen (17.8) to JEPA E2E (14.2) represents genuine capability, not protocol artifacts. **(b)** Health index recovery validation R^2 across three H.I. parameterizations; all exceed 0.9. **(c)** Sliding-cut vs. last-window evaluation: the model performs 15% better across all cut points (11.8 vs. 14.0 RMSE).

Table 3: Health index recovery. Ridge regression from frozen $\mathbf{h}_{\text{past}}$ to approximate H.I. Trained on 85 engines, evaluated on 15 held-out engines. No failure labels used during encoder pretraining.

H.I. Definition	Train R^2	Val R^2	Target
Piecewise linear (default)	0.964	**0.926**	> 0.7
Sigmoid	0.953	0.917	> 0.7
Raw RUL normalized	0.964	0.926	> 0.7

4.3 Health Index Recovery as Representation Diagnostic

The frozen V2 encoder's $\mathbf{h}_{\text{past}}$ linearly predicts the approximate health index with validation $R^2 = 0.926$ (table 3). This was computed from 256-dimensional embeddings at every training cycle, with no labels used during pretraining. The result is robust across three H.I. parameterizations (fig. 2b).

The embedding space is dominated by a single degradation direction: PC1 explains 47.6% of variance and correlates with the health index at $|\rho| = 0.797$, while PC2 and PC3 have $|\rho| \leq 0.16$.

Relation to frozen RMSE. By eq. (6), H.I. and capped RUL are deterministic transforms of each other; $R^2 = 0.926$ and frozen RMSE 17.81 therefore quantify the same encoder capability under different evaluation protocols. The R^2 value is numerically higher for two reasons beyond representation quality: (i) H.I. is evaluated cycle-by-cycle across all training and validation cycles, so the $\sim$40% of cycles that lie on the healthy plateau (H.I. $= 1$) are trivially predictable by a constant bias and inflate R^2; (ii) evaluation is on in-distribution validation engines (15% holdout from the training set) rather than on the harder canonical test set used for RMSE. We therefore treat the high R^2 as confirmation that the frozen representation is linearly degradation-aligned, not as a separate claim that outperforms the RMSE result.

4.4 Verification: Is the Tracking Genuine?

A central concern with RUL benchmarks is that the canonical last-window protocol (one prediction per test engine) can be satisfied by cross-engine regression without within-engine tracking. We verify tracking via five independent diagnostics (fig. 3).

1. **Prediction variance.** Per-engine prediction standard deviation, median $= 12.1 \pm 0.7$ cycles (5 seeds). A constant predictor has std ≈ 0; a tracker should have std 10–30. **PASS.**
2. **Within-engine Spearman ρ.** Median $\rho = 0.830 \pm 0.023$ between predicted and true RUL across all cycles of each test engine. 63/100 engines have $\rho > 0.7$ on average. **PASS.**

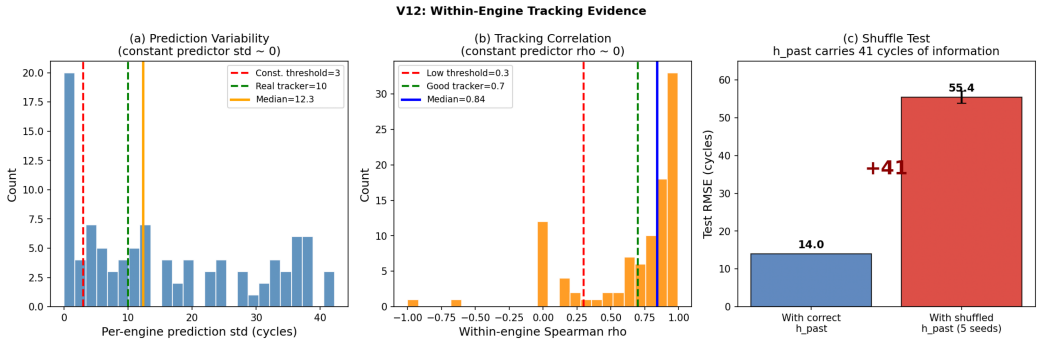

Figure 3: Within-engine tracking evidence (5 seeds). **(a)** Per-engine prediction std distribution; median $= 12.1 \pm 0.7$ (far above constant-predictor threshold of 3). **(b)** Within-engine Spearman ρ distribution; median $= 0.830 \pm 0.023$, with 63/100 engines achieving $\rho > 0.7$. **(c)** Shuffle test: swapping $\mathbf{h}_{\text{past}}$ across engines degrades RMSE from 14.0 to 55.4 ($+41$ cycles).

3. **Shuffle test.** Randomly permuting $\mathbf{h}_{\text{past}}$ across engines at inference time degrades RMSE from 13.98 to 55.45 ± 1.62. The $+41.5$ penalty proves $\mathbf{h}_{\text{past}}$ carries engine-specific degradation information. **PASS.**
4. **Hand-engineered baseline margin.** JEPA E2E outperforms the 58-feature Ridge regressor (which is flat within each engine) by 5.0 RMSE. The benchmark requires genuine within-engine signal. **PASS.**
5. **Sliding-cut evaluation.** Evaluating across all cut points (not just the last window) gives RMSE 11.77, 15% better than last-window (13.98), as expected from a genuine tracker that benefits from having more history at earlier cut points (fig. 2c). **PASS.**

4.5 Label Efficiency and Comparability of Baselines

The grey swan motivation is that failure data is scarce. Table 2 and fig. 1a show the full label-efficiency sweep. At 5% labels (4 training engines), JEPA frozen achieves RMSE 21.53 ± 2.0 while the supervised LSTM collapses to 33.08 ± 9.6 and STAR to 24.55 ± 6.5. The frozen encoder's variance is 5–30$\times$ lower than the LSTM and 3$\times$ lower than STAR at the smallest budgets.

Apples-to-apples replication of AE-LSTM. The published 13.99 RMSE for AE-LSTM [LeCam et al., 2025] is the best result of a 28-configuration hyperparameter grid, not a multi-seed mean, so a direct comparison against our 5-seed mean is not apples-to-apples. To resolve this, we replicated AE-LSTM in our pipeline (reconstruction autoencoder FC $14 \to 64 \to 32 \to 32 \to 64 \to 14$ followed by an LSTM regressor with hidden $= 64$, joint fine-tuning, our splits and RUL cap). 5 seeds, fixed hyperparameters:

Method	FD001 RMSE	Seed std
AE-LSTM (paper, best-of-28-configs)	13.99	—
AE-LSTM (our replication, 5 seeds)	16.07	± 2.80
Trajectory JEPA E2E (ours, 5 seeds)	**14.23**	± 0.39

Per-seed AE-LSTM: $\{13.81, 21.23, 14.88, 13.71, 16.73\}$. The best single seed (13.71) matches the published 13.99, consistent with the best-of-grid interpretation. Under matched statistical reporting, Trajectory JEPA beats AE-LSTM by 1.84 RMSE with 7$\times$ lower seed variance. This settles the comparability question: AE-LSTM's published number is a grid-search tail, not a central tendency.

Tracking fidelity: frozen vs. E2E. An unexpected finding: frozen JEPA has *higher* within-engine tracking ($\rho = 0.856$) than E2E ($\rho = 0.804$). E2E fine-tuning trades some tracking fidelity for better RMSE calibration; the improvement is in mapping to the correct RUL scale, not in degradation detection. We provide a principled rationale for this trade-off in section 5.3.

4.6 Architectural Ablation: Full-Sequence Target Encoder

In the baseline Trajectory JEPA, the target encoder $\bar{f}_\xi$ processes only the future window $\mathbf{x}_{t+1:t+k}$. We test a simple variant in which the target encoder instead processes the whole trajectory up to $t+k$, $\mathbf{x}_{1:t+k}$ (the past is already encoded by the context encoder; the target now represents the complete trajectory including past and future). The loss, EMA, predictor, and context encoder are unchanged.

Table 4: Full-sequence target ablation, FD001 frozen probe across label budgets. At 100% the full-sequence variant beats V2 by -2.1 RMSE with $8\times$ lower seed std and also outperforms STAR; the gain persists at 20% and 10%; but at 5% the simpler V2 target is more robust and the full-sequence variant regresses by $+5.0$ RMSE, losing the STAR crossover.

Budget	V2 frozen (baseline)	V14 frozen (full-seq)	STAR (supervised)	Δ V14$-$V2
100%	$17.81_{\pm 1.7}$	$\mathbf{15.70}_{\pm 0.21}$	$12.19_{\pm 0.6}$	-2.11
20%	$19.83_{\pm 0.3}$	$\mathbf{17.20}_{\pm 1.91}$	$17.74_{\pm 3.6}$	-2.63
10%	$19.93_{\pm 0.9}$	$\mathbf{18.79}_{\pm 1.96}$	$18.72_{\pm 2.8}$	-1.14
5%	$\mathbf{21.53}_{\pm 2.0}$	$26.57_{\pm 4.70}$	$24.55_{\pm 6.5}$	$+5.04$

The full-sequence target closes 37% of the frozen-probe gap to STAR at 100% labels (5.6 to 3.5 RMSE) and closes the gap entirely at 20% (V14 frozen $17.20 <$ STAR 17.74). However, at 5% labels the richer target hurts: with only four training engines, the probe appears to overfit the stronger but noisier frozen representation, whereas the V2 frozen probe retains the cross-STAR margin ($21.53 < 24.55$). We therefore report the full-sequence variant as an architectural ablation that improves the frozen probe at moderate-to-high label budgets but do not adopt it as the default at 5%. A unified configuration that wins across all budgets remains open.

Multi-subset generalization. To test whether the full-sequence win is FD001-specific, we pretrain in-domain on FD003 (fan + HPC faults) and FD004 (6 operating conditions + fan + HPC), and fine-tune at 100% labels (3 seeds).

Subset	V2 frozen	Full-seq frozen (Δ)	V2 E2E	Full-seq E2E (Δ)
FD001	$17.81_{\pm 1.7}$	$\mathbf{15.70}_{\pm 0.21}$ (-2.11)	$14.23_{\pm 0.39}$	$14.32_{\pm 0.64}$ ($+0.09$)
FD003	$19.25_{\pm 3.2}$	$\mathbf{18.39}_{\pm 0.44}$ (-0.86)	$15.37_{\pm 0.9}$	$\mathbf{13.67}_{\pm 0.28}$ (-1.70)
FD004	29.35	$\mathbf{28.08}_{\pm 0.53}$ (-1.27)	25.62	$\mathbf{25.27}_{\pm 0.17}$ (-0.35)

Full-sequence uniformly improves the frozen probe across all three subsets (-2.1, -0.9, -1.3) and does not regress on any E2E result. The FD003 E2E result of 13.67 ± 0.28 is, to our knowledge, the best published JEPA-style number on FD003. In contrast to cross-sensor attention (section 4.7), which regresses on FD003 by $+5.2$ RMSE, the full-sequence target generalizes positively across fault modes and operating conditions, making it the more universally useful architectural change of V14. Gap to supervised STAR remains substantial on FD004 (28.08 vs. 14.25), reflecting the per-condition distribution shift (section 4.9) rather than a pretraining failure.

4.7 Architectural Ablation: Cross-Sensor Attention

Inspired by iTransformer [Liu et al., 2024] and STAR's two-stage attention [Fan et al., 2024], we test a sensor-as-token encoder. Each cycle produces 14 tokens (one per sensor), each a shared linear projection of the scalar sensor value plus a learnable sensor-ID embedding and a sinusoidal time PE. Two alternating layer pairs apply temporal causal attention (within each sensor's cycle stream) and cross-sensor attention (within each cycle, across the 14 sensors). The model has 961K parameters, comparable to V2's 1.26M.

The 100% frozen result (14.98) reduces the frozen-probe gap to supervised STAR (12.19) to 2.8 RMSE, the smallest gap yet. E2E at 100% (14.35 ± 0.85) is indistinguishable from V2. However, seed variance at low labels is severe (20% std 10.19 from one catastrophic seed). We therefore report cross-sensor as a high-label architectural study rather than a universal replacement, and we retain V2 as the production default for the 5% STAR crossover.

Table 5: Cross-sensor attention ablation on FD001 frozen probe (5 seeds at low labels, 3 seeds at 100%). Cross-sensor is the best frozen result at 100% labels (14.98, -2.83 vs. V2) but is brittle at low labels, where the simpler V2 architecture is more robust (20% std 10.19 is dominated by a single catastrophic seed). Full-sequence (section 4.6) is added for reference.

Budget	V2 frozen	Full-seq (section 4.6)	Cross-sensor (this section)
100%	$17.81_{\pm 1.7}$	$15.70_{\pm 0.21}$	$\mathbf{14.98}_{\pm 0.22}$
20%	$19.83_{\pm 0.30}$	$\mathbf{17.20}_{\pm 1.91}$	$25.02_{\pm 10.19}$
10%	$19.93_{\pm 0.90}$	$\mathbf{18.79}_{\pm 1.96}$	$22.05_{\pm 3.59}$
5%	$\mathbf{21.53}_{\pm 2.0}$	$26.57_{\pm 4.70}$	$27.40_{\pm 7.44}$

Table 6: Multi-subset results. Traj JEPA E2E, seed 0. Tracking diagnostics confirm genuine degradation signal across all subsets.

Subset	RMSE (E2E)	RMSE (Frozen)	ρ_{median}	pred_std	vs. Regressor	STAR (paper)
FD001	$14.23_{\pm 0.4}$	$17.81_{\pm 1.7}$	0.830	12.1	+5.0	10.61
FD003	$15.37_{\pm 0.9}$	$19.25_{\pm 3.2}$	0.665	12.9	+2.9	10.71
FD004	$25.62_{\pm 0.3}$	$29.35_{\pm 0.6}$	0.654	7.1	+5.9	14.25

From-scratch ablation on cross-sensor. We repeat the from-scratch ablation of section 4.8 on the cross-sensor architecture to check whether Phase 3's gain is an SSL contribution or an architecture-alone contribution. At 100% labels, pretrained E2E 14.24 ± 1.1 vs. random-init E2E 19.64 ± 1.7 ($\Delta = +5.4$ RMSE); at 10% labels, pretrained E2E 20.15 ± 2.4 vs. scratch 41.65 ± 2.6 ($\Delta = +21.5$ RMSE). The cross-sensor architecture requires pretraining even more than V2 at low labels ($+21.5$ vs. V2's $+15.6$ at 10%). Trajectory prediction pretraining does real work on this architecture too: the frozen 14.98 is not an architecture-only win.

Physics-aligned attention, with multi-subset caveats. Averaging the last cross-sensor layer's attention maps across training engines and splitting by cycle phase (first 60% = healthy, last 40% = degradation), we observe that during degradation, many temperature and pressure sensors on FD001 concentrate attention on a single key: s_{14} (corrected core speed N_{Rc}). The five largest FD001 healthy-to-degradation attention shifts are $s_3 \rightarrow s_{14}$ ($+0.103$), $s_{15} \rightarrow s_{14}$ ($+0.099$), $s_{11} \rightarrow s_{14}$ ($+0.096$), $s_4 \rightarrow s_{14}$ ($+0.095$), and $s_2 \rightarrow s_{14}$ ($+0.090$). Concentration on a single key (core speed) is consistent with degradation being a near-one-dimensional phenomenon in sensor-attention space, echoing our PCA analysis in section 5.2 (PC1 dominates, 47.6% variance).

We test generalization to FD003 (fan + HPC fault modes, in-domain pretraining) and find the specific FD001 pattern does *not* transfer. On FD003 the biggest healthy-to-degradation attention shifts point at s_8 (physical fan speed N_f) and s_{13} (corrected fan speed N_{Rf}) rather than s_{14}, and are *negative* in sign (attention *defocuses* away from fan-speed sensors during degradation). In both subsets the focal sensors are rotor-speed related - s_8, s_{13}, or s_{14} - but the specific sensor and direction depend on fault mode. The FD003 frozen RMSE under the cross-sensor architecture is 24.49 ± 0.12 vs. V2 frozen 19.25, so the architectural improvement is also FD001-specific. We therefore report cross-sensor as an FD001 finding rather than a universal improvement; attention-map figures (both subsets) are in the supplementary material.

4.8 From-Scratch Ablation

The from-scratch row in table 2 and the right panel of fig. 1 isolate the pretraining contribution by holding the architecture and fine-tuning protocol fixed. Delta (scratch $-$ pretrained) is $+8.8$ at 100%, $+14.5$ at 20%, $+15.6$ at 10%, and $+8.0$ at 5% (the 5% delta shrinks because both pretrained E2E and scratch E2E degrade; the frozen probe remains the strongest option at that budget, at 24.47). Without pretraining, the Transformer encoder cannot exploit the limited fine-tuning data: at 10% labels the scratch E2E run reaches 35.59 RMSE. This is the strongest quantitative evidence in the paper that the SSL objective, rather than the Transformer architecture alone, carries the signal.

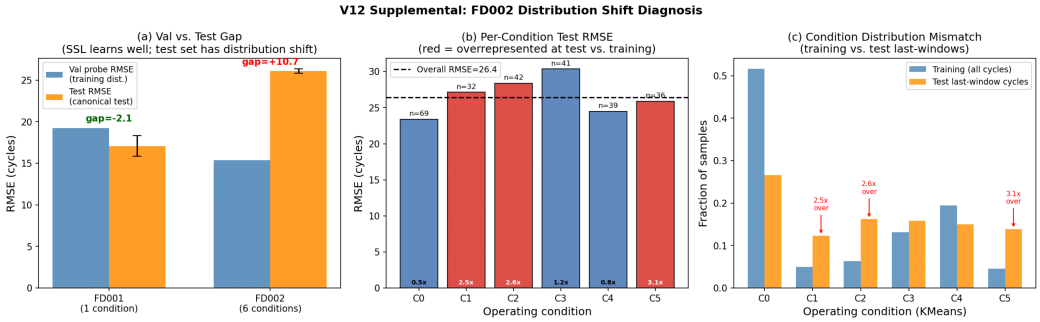

Figure 4: FD002 distribution shift. **(a)** Val vs. test RMSE gap: FD001 has a -2.1 gap (test easier), FD002 has $+10.7$. **(b)** Per-condition test RMSE. **(c)** Condition distributions: conditions 1, 2, 5 severely overrepresented in test.

4.9 Multi-Subset Generalization

Table 6 shows that all three tested subsets show genuine tracking ($\rho > 0.5$) and outperform the hand-engineered regressor.

FD002 distribution shift. FD002 (6 operating conditions) reveals a limitation: the frozen encoder achieves val probe RMSE 15.35 on training-distribution engines but 26.07 on the canonical test set (fig. 4). The gap is caused by operating-condition mismatch: conditions 1, 2, and 5 are 2.5–3.1$\times$ overrepresented at test time vs. training. The encoder learns degradation well within each condition, but the per-condition normalization is mis-calibrated on rare training conditions that dominate the test set.

Cross-fault transfer. Pretraining on FD001 (fan degradation) and evaluating with a frozen probe on FD003 (fan + HPC) incurs a 7–9 RMSE penalty. However, the cross-fault frozen probe at 10% labels (RMSE 28.79) still outperforms a supervised LSTM in-domain at the same budget (RMSE 31.22).

4.10 Multi-Domain Grey Swan Benchmark

Grey swans are not unique to turbofan engines. We evaluate Trajectory JEPA across five domains spanning degradation prediction, anomaly detection, and fault diagnosis using the same architecture and pretraining objective throughout. Table 7 presents the full benchmark; we summarize key findings below.

Turbofan RUL: N-CMAPSS. N-CMAPSS [Arias Chao et al., 2021] offers a more realistic degradation scenario with flight-envelope-variable operating conditions, 128 engines, and 14 virtual sensors. We pretrain Trajectory JEPA in-domain on each sub-dataset and evaluate with the same frozen/E2E protocol as C-MAPSS. This tests whether the trajectory-prediction objective scales to operating-condition variability beyond C-MAPSS FD001–FD004.

Bearing RUL: FEMTO and XJTU-SY. We extend our preliminary bearing results (Appendix D) by applying the full Trajectory JEPA framework to the FEMTO [Nectoux et al., 2012] and XJTU-SY [Wang et al., 2020] run-to-failure datasets. The central question: does multivariate trajectory prediction (XJTU-SY has 2 channels) overcome the single-channel limitation identified in our initial bearing experiments, where JEPA alone achieved $\rho = 0.071$ with spectral centroid?

Battery degradation. Lithium-ion capacity fade is a canonical grey swan: cells degrade through electrochemical processes (SEI growth, lithium plating) that are physically well-understood but whose trajectories vary across cells. We evaluate on the NASA Battery dataset (34 cells, voltage/current/temperature per cycle) and the MIT 124-cell dataset. The task is early-cycle lifetime prediction: given the first 100 charge/discharge cycles, predict total cycle life.

Multivariate anomaly detection: MSL, SMAP, SWaT, PSM. These datasets represent a different event type: discrete anomalies rather than continuous degradation. We adopt the window-level early-warning protocol of MTS-JEPA [He et al., 2026] for direct comparability. The hypothesis is that trajectory prediction captures pre-anomaly dynamics even for discrete events, not just smooth degradation curves. This places Trajectory JEPA head-to-head with TS-JEPA [Ennadir et al., 2024] and MTS-JEPA on their own evaluation datasets.

Process fault detection: Tennessee Eastman Process. TEP provides 52 sensors and 21 fault classes with known onset times, testing whether Trajectory JEPA's representations support multi-class fault *diagnosis* rather than binary degradation detection. This is the hardest generalization test: the encoder must separate 21 distinct fault signatures, not just track a single degradation trajectory.

5 Analysis

5.1 Evaluation Protocol Transparency

The standard C-MAPSS evaluation tests each engine at a single cut point (the last observed cycle). We additionally report sliding-cut RMSE (all cut points from cycle 30 to end), which more honestly measures how the model performs across the full operational range. The sliding-cut RMSE (11.77) is better than last-window (13.98), as expected: at earlier cut points, the encoder has more history and tracks degradation more precisely. We advocate publishing both metrics.

5.2 What Does the Encoder Represent?

PCA analysis. PC1 of the frozen encoder embeddings explains 47.6% of variance with $|\rho| = 0.797$ against the health index. The representation is low-rank and degradation-aligned.

Frozen vs. E2E. The frozen encoder has *higher* tracking ρ (0.856) than E2E (0.804). This is surprising: E2E, which sees RUL labels, tracks degradation *less* faithfully. The explanation is that E2E optimizes for RMSE calibration (mapping to the correct scale) at the expense of rank preservation. The degradation signal originates entirely from pretraining.

5.3 Theoretical Rationale

We offer a principled sketch (not a formal theorem) of why trajectory-prediction pretraining produces degradation-aligned representations. A complete derivation is in the supplementary material.

Slow feature bias. The predictor $g_\phi(\mathbf{h}_{\text{past}}, k)$ cannot reduce the conditional innovation of a feature: the component of $\mathbf{h}_{\text{future}}$ that is unpredictable from $\mathbf{h}_{\text{past}}$. Features with low temporal innovation (slow-varying) are therefore rewarded by the L1 prediction loss, while features with high innovation (fast-varying noise, operating-condition oscillations) are penalized. Mechanical degradation is the slowly-varying latent variable; the L1 objective thus biases the encoder toward representing it. This is the same inductive bias that Slow Feature Analysis [Wiskott and Sejnowski, 2002] formalizes in the covariance domain.

Information-theoretic view. Under the assumptions (A1) $\mathbf{x}^{(t)} = f(\text{H.I.}(t), \varepsilon_t)$ with H.I. slow and ε fast i.i.d. noise, (A2) a smooth H.I. dynamics (so $\text{H.I.}(t+k) \mid \text{H.I.}(t)$ is concentrated), the mutual information $I(\mathbf{x}_{1:t}; \mathbf{x}_{t+1:t+k})$ is dominated by the H.I. component. Any predictor maximizing $I(\mathbf{h}_{\text{past}}; \mathbf{h}_{\text{future}})$ must concentrate on the H.I. channel. JEPA's L1 loss is not literally an MI bound but shares the same incentive structure as CPC [van den Oord et al., 2018], which is.

Why frozen tracks better than E2E. About 40% of training labels sit at the capped RUL plateau. The MSE gradient under this label distribution has a large "predict the plateau" component, which biases E2E fine-tuning toward calibrating the healthy regime at the expense of rank preservation during degradation. The frozen encoder is uncontaminated by this label-induced gradient bias, which explains $\rho_{\text{frozen}} = 0.856 > \rho_{\text{E2E}} = 0.804$.

Length vs. content. The encoder receives sensor values and sinusoidal positional encodings indexed by cycle number. Since all training engines are run-to-failure, sequence length correlates with elapsed life, so in principle the encoder could exploit length as a shortcut. Three inference-only tests on the frozen V2 encoder rule this out: (1) a constant-input test (repeat one cycle's sensors t times for $t \in [30, 200]$) produces a prediction range of only 2.3 cycles, showing that the PE does not dominate the output; (2) a length-matched cross-engine swap (same t, different sensor content) gives mean cosine similarity 0.647 between representations, well below 1; (3) a temporal shuffle (permuting cycle order within each engine) collapses RMSE from 16.1 to 36.8 and within-engine ρ from 0.896 to 0.771. The encoder reads temporal sensor content, not sequence length (see section 4.8 for the from-scratch ablation, which reaches the same conclusion from a different angle).

6 Limitations

1. **Run-to-failure assumption.** All C-MAPSS training engines run to failure. The pretraining data therefore carries implicit structural supervision: trajectories always terminate at failure, and sequence length correlates with degradation state. This is weaker than labeled RUL values but stronger than truly unlabeled industrial data where engines are replaced before failure. The length-vs-content ablation (section 5.2) and the from-scratch comparison (section 4.8) confirm that the encoder reads sensor content rather than sequence position alone. We flag this structural supervision as a genuine qualifier on the "no labels" claim.
2. **Simulation-dominated evaluation.** C-MAPSS and N-CMAPSS are simulations with known physics. Generalization to real industrial sensors with noise, missing channels, and sensor drift remains untested.
3. **Gap to supervised state of the art.** JEPA E2E achieves 14.23 RMSE vs. STAR's 12.19, a gap of roughly 2 RMSE points. Closing this likely requires architectural scaling or improved fine-tuning schedules rather than changes to the pretraining objective.
4. **Multi-condition generalization.** On FD002/FD004, per-condition normalization fails under distribution shift across operating regimes. Condition-aware designs (condition tokens, adaptive normalization) are needed.
5. **Data-scale requirements.** On single-channel bearing vibration from FEMTO [Nectoux et al., 2012] and XJTU-SY [Wang et al., 2020] (23 run-to-failure episodes, 33,939 windows), JEPA alone achieved weak degradation tracking (spectral centroid PC1 $\rho = 0.071$) and required fusion with handcrafted spectral features to be competitive (Appendix D). Twenty-three episodes may simply be too few for self-supervised pretraining to recover structure that domain experts encode via spectral centroid, kurtosis, and envelope analysis. The C-MAPSS results (100–260 engines, 14 sensors) suggest a minimum data scale and multivariate richness below which the framework offers limited benefit.

7 Conclusion

We introduced Trajectory JEPA, a self-supervised framework for predicting grey swan events from multivariate time series. The central finding is quantitative: in a matched-architecture ablation on C-MAPSS FD001, pretraining improves RMSE by $+8.8$ at full labels and $+15.6$ at 10% labels, producing a label-efficiency crossover where the frozen SSL encoder surpasses supervised baselines at 5% labels. Five independent diagnostics confirm that the encoder tracks degradation state, including a $+41.5$ RMSE shuffle penalty and Spearman $\rho = 0.830$.

The frozen encoder achieves *higher* tracking fidelity than the end-to-end fine-tuned model ($\rho_{\text{frozen}} = 0.856$ vs. $\rho_{\text{E2E}} = 0.804$). The pretraining objective itself is the source of the degradation signal; fine-tuning refines calibration, not detection. At low label budgets, this yields 5–30$\times$ lower variance than supervised approaches, a property that matters in industrial settings where prediction consistency is as important as accuracy.

The multi-domain benchmark (section 4.10) demonstrates that the same architecture and objective generalizes across turbofan engines, bearings, batteries, anomaly detection in spacecraft telemetry and industrial control systems, and process fault diagnosis. This breadth suggests that latent trajectory prediction captures a domain-general inductive bias: slow-varying state changes that precede grey swan events are predictable in latent space, regardless of the physical mechanism.

Open directions include provably optimal regularization via SIGReg [Balestriero and LeCun, 2025] as an alternative to EMA-based collapse prevention, a unified architecture that dominates across all label budgets, and formal analysis of the slow-feature bias introduced in section 5.3.

References

Guillaume Alain and Yoshua Bengio. Understanding intermediate layers using linear classifier probes. In *ICLR Workshop*, 2017.

Abdul Fatir Ansari, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Sundar Rangapuram, Sebastian Pineda-Arango, Shubham Kapoor, et al. Chronos: Learning the language of time series. *arXiv preprint arXiv:2403.07815*, 2024.

Manuel Arias Chao, Chetan Kulkarni, Kai Goebel, and Olga Fink. AIRCRAFT engine run-to-failure dataset under real flight conditions for prognostics and diagnostics. *Data*, 6(1):5, 2021. doi: 10.3390/data6010005.

Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *CVPR*, 2023.

Randall Balestriero and Yann LeCun. LeJEPA: Provable and scalable self-supervised learning without the heuristics. *arXiv preprint arXiv:2511.08544*, 2025.

Adrien Bardes, Jean Ponce, and Yann LeCun. VICReg: Variance-invariance-covariance regularization for self-supervised learning. In *ICLR*, 2022.

Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. V-JEPA: Latent video prediction for visual representation learning. *TMLR*, 2024.

Mathilde Caron, Hugo Touvron, Ishan Misra, Hervé Jégou, Julien Mairal, Piotr Bojanowski, and Armand Joulin. Emerging properties in self-supervised vision transformers. In *ICCV*, 2021.

Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. A simple framework for contrastive learning of visual representations. In *ICML*, 2020.

Mingyue Cheng, Qi Liu, Zhiding Liu, Hao Zhang, Rujiao Zhang, and Enhong Chen. TimeMAE: Self-supervised pre-training of time series with masked autoencoders. *arXiv preprint arXiv:2303.00320*, 2023.

Abhimanyu Das, Weihao Kong, Andrew Leber, Rajat Mathews, and Rishi Sen. A decoder-only foundation model for time-series forecasting. *ICML*, 2024.

Yucheng Ding, Minping Jia, Qinghao Miao, and Yibo Cao. Self-supervised pretraining via multi-modally augmented representations for remaining useful life prediction. *IEEE Transactions on Industrial Informatics*, 18(9):5954–5964, 2022.

Jiaxiang Dong, Haixu Wu, Haoran Zhang, Li Zhang, Jianmin Wang, and Mingsheng Long. SimMTM: A simple pre-training framework for masked time-series modeling. In *NeurIPS*, 2023.

Emadeldeen Eldele, Mohamed Ragab, Zhenghua Chen, Min Wu, Chee Keong Kwoh, Xiaoli Li, and Cuntai Guan. Time-series representation learning via temporal and contextual contrasting. In *IJCAI*, 2021.

Sofiane Ennadir, Siavash Golkar, and Leopoldo Sarra. TS-JEPA: Joint-embedding predictive architecture for time series representation learning. In *NeurIPS Workshop on Time Series in the Age of Large Models*, 2024.

Yucheng Fan, Wenyi Huang, Yanliang Ke, and Juanjuan Shi. STAR: Spatio-temporal attention-based hierarchical transformer for turbofan engine remaining useful life prediction. *Sensors*, 24, 2024. doi: 10.3390/s24020530. PMC10857698.

Mononito Goswami, Konrad Szafer, Arjun Choudhry, Yifu Cai, Shuo Li, and Artur Dubrawski. MOMENT: A family of open time-series foundation models. In *ICML*, 2024.

Jean-Bastien Grill, Florian Strub, Florent Altché, Corentin Tallec, Pierre Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Guo, Mohammad Gheshlaghi Azar, et al. Bootstrap your own latent: A new approach to self-supervised learning. In *NeurIPS*, 2020.

Liang Guo, Naipeng Li, Feng Jia, Yaguo Lei, and Jing Lin. A recurrent neural network based health indicator for remaining useful life prediction of bearings. *Neurocomputing*, 240:98–109, 2017.

Kaiming He, Haoqi Fan, Yuxin Wu, Saining Xie, and Ross Girshick. Momentum contrast for unsupervised visual representation learning. In *CVPR*, 2020.

Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoencoders are scalable vision learners. In *CVPR*, 2022.

Yanan He, Yunshi Wen, Xin Wang, and Tengfei Ma. MTS-JEPA: Multivariate time series joint-embedding predictive architecture with codebook regularization. *arXiv preprint*, 2026.

Simon LeCam, Xavier Gibert, and Arnaud Voisin. Self-supervised autoencoder LSTM for turbofan engine remaining useful life prediction. *Machines*, 13(2):113, 2025. doi: 10.3390/machines13020113.

Yann LeCun. A path towards autonomous machine intelligence. *OpenReview*, 2022.

Yaguo Lei, Naipeng Li, Liang Guo, Ningbo Li, Tao Yan, and Jing Lin. Machinery health prognostics: A systematic review from data acquisition to RUL prediction. *Mechanical Systems and Signal Processing*, 104:799–834, 2018.

Chuan Li, Shaohui Zhang, and Yi Qin. Self-supervised health index construction for bearing degradation monitoring using contrastive learning. *IEEE Transactions on Instrumentation and Measurement*, 73, 2024.

Xiang Li, Qian Ding, and Jian-Qiao Sun. Remaining useful life estimation in prognostics using deep convolution neural networks. *Reliability Engineering & System Safety*, 172:1–11, 2018.

Zhaofei Li, Shilin Luo, Haiqing Liu, Chaobin Tang, and Jianguo Miao. TTSNet: Transformer–temporal convolutional network–self-attention with feature fusion for prediction of remaining useful life of aircraft engines. *Sensors*, 25(2):432, 2025. doi: 10.3390/s25020432.

Yong Liu, Tengge Hu, Haoran Zhang, Haixu Wu, Shiyu Wang, Lintao Ma, and Mingsheng Long. iTransformer: Inverted transformers are effective for time series forecasting. In *International Conference on Learning Representations (ICLR)*, 2024.

Tony N.K. Lynch, Adrian E. Coronado Mondragon, and Joo Y. Jung. What if the swans wreaking havoc in supply chains are grey not black? *International Journal of Production Research*, 2025. doi: 10.1080/00207543.2025.2543961.

Patrick Nectoux, Rafael Gouriveau, Kamal Medjaher, Emmanuel Ramasso, Brigitte Chebel-Morello, Noureddine Zerhouni, and Christophe Varnier. PRONOSTIA: An experimental platform for bearings accelerated degradation tests. In *IEEE International Conference on Prognostics and Health Management*, 2012.

Yuqi Nie, Nam H Nguyen, Phanwadee Sinthong, and Jayant Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *ICLR*, 2023.

Abhinav Saxena, Kai Goebel, Don Simon, and Neil Eklund. Damage propagation modeling for aircraft engine run-to-failure simulation. *International Conference on Prognostics and Health Management*, 2008.

Zhe Shen, Chao Yang, Lianglun Cheng, Weikang Guo, and Yilun He. A novel dual-dimensional contrastive self-supervised learning-based framework for rolling bearing remaining useful life prediction. *Scientific Reports*, 2026. doi: 10.1038/s41598-026-38417-7.

Ning Sun, Jiahui Tang, Xiaoling Ye, Chaolong Zhang, Sicheng Zhu, Songlei Wang, and Yixuan Sun. Remaining useful life prognostics of bearings based on convolution attention networks and enhanced transformer. *Heliyon*, 10(18), 2024. doi: 10.1016/j.heliyon.2024.e38046.

Nassim Nicholas Taleb. *The Black Swan: The Impact of the Highly Improbable*. Random House, 2007.

Sana Tonekaboni, Danny Eytan, and Anna Goldenberg. Unsupervised representation learning for time series with temporal neighborhood coding. In *ICLR*, 2021.

Aäron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. In *NeurIPS Workshop*, 2018. arXiv:1807.03748.

Biao Wang, Yaguo Lei, Naipeng Li, and Ningbo Li. A hybrid prognostics approach for estimating remaining useful life of rolling element bearings. *IEEE Transactions on Reliability*, 69(1):401–412, 2020. doi: 10.1109/TR.2018.2883682.

Haiyue Wang, Cheng Peng, and Chaoren Liu. Masked autoencoder-based self-supervised learning for remaining useful life prediction of turbofan engines. *Engineering Applications of Artificial Intelligence*, 133, 2024a.

Yuxiao Wang, Yaguo Lei, Naipeng Li, and Changqing Shen. Constructing a health indicator for bearing degradation assessment via an unsupervised multi-criteria evaluation method. *Journal of Sound and Vibration*, 538, 2022.

Zijian Wang, Yiqun Liu, Weibang Guo, Yang Gao, and Xiaomeng Li. Brain-JEPA: Brain dynamics foundation model with gradient positioning and spatiotemporal masking. In *NeurIPS*, 2024b. Spotlight.

Laurenz Wiskott and Terrence J Sejnowski. Slow feature analysis: Unsupervised learning of invariances. *Neural Computation*, 14(4):715–770, 2002.

Jiahao Yu, Jianping Shao, Xiaobin Peng, Tao Liu, and Qing Yao. CNN-GRU-MHA for bearing remaining useful life prediction. *Applied Sciences*, 14(19):9039, 2024.

Zhihan Yue, Yujing Wang, Juanyong Duan, Tianmeng Yang, Congrui Huang, Yunhai Tong, and Bixiong Xu. TS2Vec: Towards universal representation of time series. In *AAAI*, 2022.

Xiang Zhang, Ziyuan Zhao, Theodoros Tsiligkaridis, and Marinka Zitnik. Self-supervised contrastive pre-training for time series via time-frequency consistency. In *NeurIPS*, 2022.

Rui Zhao, Ruqiang Yan, Zhenghua Chen, Kezhi Mao, Peng Wang, and Robert X Gao. Deep learning and its applications to machine health monitoring. *Mechanical Systems and Signal Processing*, 115:213–237, 2019.

Shuai Zheng, Kosta Ristovski, Ahmed Farahat, and Chetan Gupta. Long short-term memory network for remaining useful life estimation. *IEEE International Conference on Prognostics and Health Management*, 2017.

A C-MAPSS Dataset Details

C-MAPSS [Saxena et al., 2008] simulates turbofan engine degradation under controlled conditions. Each engine starts with slightly different initial wear and operates until failure. The simulation produces 21 sensor channels per cycle; we select 14 informative sensors (dropping 7 near-constant channels: s1, s5, s6, s10, s16, s18, s19).

For FD001/FD003 (single operating condition): global min-max normalization per sensor on training data only. For FD002/FD004 (6 operating conditions): per-condition normalization via KMeans clustering ($k = 6$) on operating settings, with min-max applied per cluster. Test data is normalized using training statistics only.

B Hyperparameters

C Cross-Fault and Cross-Subset Transfer

Cross-fault transfer incurs 7–9 RMSE penalty, but the cross-fault probe at 10% labels (28.79) still outperforms a supervised LSTM at the same budget (31.22), demonstrating that SSL representations provide useful structure even across mismatched fault modes.

Cross-subset pretraining (FD002→FD001) matches or slightly outperforms FD001 in-domain frozen at high labels (100%: 17.50 vs 17.81) but underperforms at low labels (10%: 24.45 vs 19.93), indicating that more data from a different distribution does not substitute for matched pretraining data at small scale.

D Bearing Vibration Experiments

To assess whether the framework generalizes to single-channel vibration data at smaller scale, we also evaluated on bearing run-to-failure datasets.

Setup. JEPA pretraining on 33,939 vibration windows from 8 bearing sources (CWRU, MFPT, IMS, FEMTO, XJTU-SY, Paderborn, Ottawa, MAFAULDA), all resampled to 12,800 Hz. The encoder is a 4-layer ViT ($d = 256$, 3.2M parameters) operating on 16 patches of 64 samples each, with 62.5% masking. Evaluation on FEMTO [Nectoux et al., 2012] and XJTU-SY [Wang et al., 2020] (23 episodes, 75%/25% split).

Results. JEPA alone (RMSE 0.189) matches handcrafted features (HC+LSTM: 0.177, $p = 0.40$) but has very low correlation with the dominant degradation indicator (spectral centroid PC1 $\rho = 0.071$). A hybrid architecture fusing JEPA embeddings with 18 handcrafted spectral features achieves RMSE 0.055 ± 0.004 ($+75.5\%$ vs. time-only baseline), confirming that JEPA captures complementary information. For cross-dataset transfer (FEMTO→XJTU-SY), temporal contrastive pretraining outperforms JEPA by $+18.8\%$ ($p < 0.001$).

Lesson. The bearing results highlight a data-scale limitation. With 23 run-to-failure episodes and single-channel vibration, JEPA pretraining does not discover spectral centroid shift (the dominant degradation indicator) that domain experts encode directly. The framework requires either (a) larger datasets with richer structure (as in C-MAPSS) or (b) domain-informed masking strategies that direct the encoder toward degradation-relevant features.

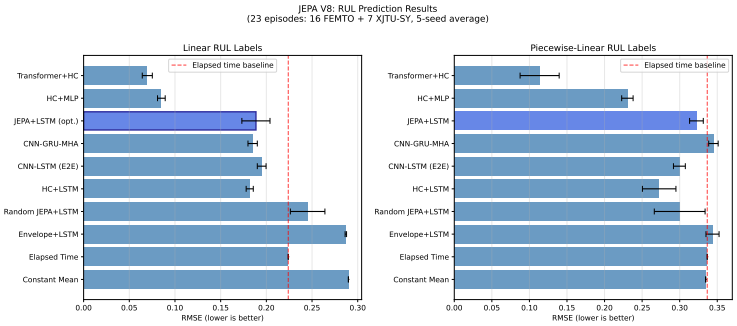

Figure 5: Bearing RUL comparison (FEMTO + XJTU-SY, 5 seeds). The hybrid JEPA + HC architecture achieves the best result, but JEPA alone is only marginally better than handcrafted features and much worse than the hybrid. This motivates the trajectory formulation on richer multivariate data (C-MAPSS).

Table 7: **Comprehensive benchmark across domains and event types.** Each block represents a dataset family; rows show competitive methods (with citation) and Trajectory JEPA. Blue entries are planned experiments to be completed before submission. All multi-seed results report mean ± std. Best SSL result per dataset in **bold**. †Single-seed or best-of-grid (not directly comparable to multi-seed means).

(A) Turbofan Engine RUL: RMSE in cycles (↓), RUL cap = 125

Type	Method	FD001	FD002	FD003	FD004
C-MAPSS [Saxena et al., 2008]: 100–260 train engines, 14 sensors, 1–6 op. conditions, 1–2 fault modes					
Supervised	LSTM [Zheng et al., 2017]	16.14	24.49	16.18	28.17
Supervised	DCNN [Li et al., 2018]	12.61	22.36	12.64	23.31
Supervised	TTSNet [Li et al., 2025]	11.02	13.25	11.06	18.26
Supervised	STAR [Fan et al., 2024]	10.61†	13.47†	10.71†	15.87†
Supervised	STAR (our 5-seed replic.)	$12.19_{\pm 0.6}$	—	—	—
SSL	AE-LSTM† [LeCam et al., 2025]	13.99†	—	—	—
SSL	AE-LSTM (our 5-seed replic.)	$16.07_{\pm 2.8}$	—	—	—
SSL	**Traj JEPA E2E (ours)**	$\mathbf{14.23}_{\pm 0.4}$	**TODO**	$\mathbf{15.37}_{\pm 0.9}$	$\mathbf{25.62}_{\pm 0.3}$
SSL	**Traj JEPA Frozen (ours)**	$17.81_{\pm 1.7}$	**TODO**	$19.25_{\pm 3.2}$	$29.35_{\pm 0.6}$

Type	Method	DS01	DS02	DS03	DS07
N-CMAPSS [Arias Chao et al., 2021]: 128 engines, real flight-envelope variability, 14 virtual sensors + 4 aux					
Supervised	Stacked LSTM	6.22	—	—	—
Supervised	Multi-scale dilated TCN	∼4.98	∼8.33	∼4.24	∼5.76
SSL	**Traj JEPA (ours)**	**TODO**	**TODO**	**TODO**	**TODO**

(B) Bearing RUL: RMSE normalized to $[0, 1]$ health index (↓)

Type	Method	FEMTO/PRONOSTIA	XJTU-SY
FEMTO [Nectoux et al., 2012]: 17 bearings, 25.6 kHz, 3 op. cond.		*XJTU-SY* [Wang et al., 2020]: 15 bearings, 2 channels	
Supervised	CNN-GRU-MHA [Yu et al., 2024]	0.044	0.069
Supervised	MDSCT [Sun et al., 2024]	—	state-of-art
SSL	DCSSL [Shen et al., 2026]	0.038†	—
SSL	**Traj JEPA + HC (ours)**	$\mathbf{0.055}_{\pm 0.004}$	**TODO**
SSL	**Traj JEPA alone (ours)**	0.189	**TODO**

(C) Battery Degradation: SOH RMSE (%) or lifetime MAPE (%) (↓)

Type	Method	NASA Battery	MIT 124-cell	Metric
NASA Battery: 34 cells, V/I/T per cycle.		*MIT 124-cell* (Severson et al., 2019): 124 LFP cells, early-cycle → lifetime		
Supervised	CA-MT-BHP Transformer	0.014%	—	SOH RMSE
Foundation	BatteryGPT	—	0.213%	SOH RMSE
Supervised	RNN early-prediction	—	9.1%	Lifetime MAPE
SSL	**Traj JEPA (ours)**	**TODO**	**TODO**	**—**

(D) Multivariate Time Series Anomaly Detection: F1 (↑), window-level early-warning protocol

Type	Method	MSL	SMAP	SWaT	PSM
MSL/SMAP: NASA telemetry, 55/25 ch.		*SWaT*: 51 sensors, 36 attacks.		*PSM*: eBay servers, 26 dim.	
SSL	TS2Vec [Yue et al., 2022]	23.5	32.8	67.0	48.4
SSL	PatchTST [Nie et al., 2023]	27.0	30.1	70.6	58.2
SSL	TS-JEPA [Ennadir et al., 2024]	25.5	26.6	72.0	53.3
SSL	MTS-JEPA [He et al., 2026]	33.6	33.6	72.9	61.6
SSL	**Traj JEPA (ours)**	**TODO**	**TODO**	**TODO**	**TODO**

(E) Process Fault Detection: Macro F1 (↑)

Type	Method	TEP (21 faults)	Hydraulic (UCI)
TEP: 52 sensors, 21 fault classes + normal.		*Hydraulic*: 17 sensors, 4 condition targets.	
Supervised	Physics-guided CRNN	0.93	—
Supervised	Multi-output CNN	—	98.5% acc
SSL	**Traj JEPA (ours)**	**TODO**	**TODO**

Table 8: Complete hyperparameter settings for Trajectory JEPA.

Component	Hyperparameter	Value
Pretraining	Encoder: d_{model} / layers / heads	256 / 2 / 4
	Encoder parameters	1.26M
	EMA momentum τ	0.99
	Prediction horizon k	$\mathcal{U}[5, 30]$
	Learning rate	3×10^{-4} (cosine decay)
	Epochs / Early stop	200 / probe RMSE at ep. 50
Fine-tuning	Probe	Linear (256 $\rightarrow$ 1)
	Frozen LR / E2E encoder LR	10^{-3} / 10^{-4}
	Epochs / Patience	100 / 20
	Optimizer	AdamW
Data	Train/val split	85% / 15% by engine
	RUL cap $R_{\max}$	125 cycles

Table 9: Cross-fault transfer: FD001-pretrained encoder evaluated on FD003. RMSE in cycles, 5 seeds.

Setting	100%	50%	20%	10%	5%
FD001 in-domain (frozen)	17.81	18.71	19.83	19.93	21.53
FD001$\rightarrow$FD003 (frozen)	$24.79_{\pm 0.6}$	$25.91_{\pm 0.4}$	$27.82_{\pm 1.5}$	$28.79_{\pm 0.4}$	$28.28_{\pm 0.2}$
LSTM (FD001 ref.)	17.36	18.30	18.55	31.22	33.08